

CS224R Spring 2026 Homework 1

Imitation Learning

Due 4/10/2026 9pm PT

SUNet ID:

Name:

Collaborators:

By turning in this assignment, I agree by the Stanford honor code and declare that all of this is my own work.

Overview

Goals: In this assignment you will experiment with imitation learning on a custom Flappy Bird environment using action chunking. You will train behavior cloning (BC) policies with different losses regression and flow matching, and improve it with DAgger.

1. Implement and train an imitation learning agent using regression.
2. Implement and train an imitation learning agent via flow matching to learn a generative policy.
3. Implement DAgger to improve the first regression policy via iterative relabeling with a deterministic expert.

Environment

The environment is a physics-based Flappy Bird where the agent controls a target y -position (normalised to $[0, 1]$). A PD controller converts the target into thrust internally, creating momentum-based dynamics.

- **Observation** (4-D, normalised): `[dist_to_pipe, gap1.y, gap2.y, bird.y]`.
- **Action:** target y -position in $[0, 1]$.
- **Easy mode:** one gap per pipe.
- **Hard mode:** alternating single and double gap pipes.
- **Max episode length:** 1000 steps. An agent that survives all 1000 steps is considered successful.

Action chunking: Policies predict `ACTION_CHUNK=20` future target positions at once. During rollout, only the first `EXECUTE_STEPS=10` are executed before re-querying the policy (called receding horizon control).

Submission

Submitting the PDF: Make a PDF report containing the results to Problem 1, 2 and 3. And submit to gradescope on the **Homework 1 (Written Part)**.

Submitting the Code: Submit on the gradescope assignment called **Homework 1 (Programming Part)** a zip file containing:

- Your completed code files (`networks.py`, `losses.py`, `expert.py`, `dagger.py`) in the original format under the folder `hw1` with TODOs filled in.
- The results from your runs `bc_reg_easy.txt`, `bc_reg_hard.txt`, `bc_flow_hard.txt`, `dagger_hard.txt`
- The resulting folder that you will submit should look like the following:

```
.
|-- hw1/                # containing all the code
|-- bc_reg_easy.txt
|-- bc_reg_hard.txt
|-- bc_flow_hard.txt
|-- dagger_hard.txt
```

Use of AI Tools (e.g. ChatGPT, Cursor): For the sake of deeper understanding on implementing imitation learning methods, assistance from generative models to write code for this homework is prohibited. See the course website for more details.

Setup

See `installation.md` in the starter code for detailed instructions.

See `colab_instructions.md` if you prefer to run your code in google colab.

Codebase

The starter code lives in the `hw1/` directory. Your task is to fill in the functions marked with `TODO` (which raise `NotImplementedError`).

- `main.py` — training pipeline and CLI entrypoint (**read-only**)
- `visualization.py` — evaluation, video recording, policy wrappers (**read-only**)
- `flappy_bird_env.py` — Gymnasium environment (**read-only**)
- `expert.py` — Agent expert that provides the initial demos (**read-only**)
- `networks.py` — neural network architectures [**TODO**: `BCPolicy`, `FlowMatchingSchedule`]
- `losses.py` — loss functions [**TODO**: `bc_loss`, `flow_matching_loss`]
- `dagger.py` — DAgger relabeling [**TODO**: `DeterministicExpert.act`, `rollout_and_relabel`]

We recommend implementing in this order:

1. `networks.py::BCPolicy`
2. `losses.py::bc_loss`
3. `networks.py::FlowMatchingSchedule.interpolate` and `FlowMatchingSchedule.sample`
4. `losses.py::flow_matching_loss`
5. `dagger.py::DeterministicExpert.act`
6. `dagger.py::rollout_episode`
7. `dagger.py::rollout_and_relabel`

Problem 1: Behavior Cloning with Regression [2 points]

In this section, you will implement the Behavior Cloning (BC) algorithm to train your policies using the Mean Square Error (MSE) loss.

We provide a set of demonstrations of an expert playing our Stanford Flappy Bird game (see `expert.py`). This dataset $\mathcal{D}$ consists of state action pairs (s, a) .

BC entails learning a policy π that matches from the state of the flappy bird environment s , to the action to be taken by the bird a .

Today, most robot learning policies predict the set of T future actions $a_t \dots a_{t+T}$ out of which only $K < T$ are rolled out open loop (without re-querying the policy) and the rest are discarded. This technique, called action chunking, generally results in much better performance.

1. **Implementation.** Implement the following:

- `BCPolicy.__init__` and `BCPolicy.forward` in `networks.py`: a 3-layer MLP (Linear $\rightarrow$ ReLU $\rightarrow$ Linear $\rightarrow$ ReLU $\rightarrow$ Linear $\rightarrow$ Sigmoid).
- `bc_loss` in `losses.py`: MSE loss between predicted and expert actions:

$$\mathcal{L}_{\text{MSE}} = \frac{1}{N} \sum_{i=1}^N \|\hat{a}_i - a_i^*\|^2 \quad (1)$$

where $\hat{a}_i = \pi_{\theta}(s_i)$ is the predicted action and a_i^* is the expert action.

2. **Run the regression agent on easy mode** and report your results in a table (mean and standard deviation of episode length over 50 evaluation episodes) [1 point].

```
python main.py --method bc_reg --env easy
```

Solution: We expect a mean larger than 900

3. **Run the regression agent on hard mode** and report the results in a table below (no new code needed):

```
python main.py --method bc_reg --env hard
```

Solution: We expect the policy to fail something between 100 and 300 (No points assigned to this question)

4. **Explain** in 2–3 sentences the performance of the MSE regression on hard mode. Why is it performing the way it is? [1 point]

Solution: The hard mode version of the Stanford Flappy Bird has pipes with two

openings. The dataset that we train on has bimodal data, where with a random probability the agent will choose to go through the top opening or through the bottom opening. MSE is not well suited when handling multimodality of the data, and hence as a result when learning from this data MSE results in averaging out the two modes and crashes with the separation pipe.

What do we consider a correct answer: Students should realize about the multimodality of the original dataset that originates from the double pipe opening [+0.5 points]. Students should realize that MSE is not well suited to handle such multimodality [+0.5 points].

Problem 2: Flow Matching [2 points]

In this section, you will implement flow matching to train your policies instead of using an MSE loss. Flow matching learns a conditional vector field that iteratively transports noise into realistic action chunk predictions. Flow matching is similar to diffusion, but it is simpler to implement and generally exhibits similar or better performance.

Let $\mathbf{a}_t$ be an action chunk in the demonstration dataset and $\mathbf{a}_{t,0} \sim \mathcal{N}(0, I)$ be noise of the same shape. We first sample a “flow matching timestep” $\tau \sim \mathcal{U}(0, 1)$ and define the interpolation $\mathbf{a}_{t,\tau} = \tau \mathbf{a}_t + (1 - \tau) \mathbf{a}_{t,0}$. We then train a network v_θ to predict the velocity that moves $\mathbf{a}_{t,\tau}$ toward $\mathbf{a}_t$, using the flow-matching loss:

$$\mathcal{L}_{\text{FM}}(\theta) = \frac{1}{|\mathcal{D}|} \sum_{(s_t, \mathbf{a}_t) \in \mathcal{D}} \|v_\theta(s_t, \mathbf{a}_{t,\tau}, \tau) - (\mathbf{a}_t - \mathbf{a}_{t,0})\|_2^2. \quad (2)$$

At inference time, we sample initial noise $\mathbf{a}_{t,0} \sim \mathcal{N}(0, I)$ and integrate the ODE $\frac{d\mathbf{a}_{t,\tau}}{d\tau} = v_\theta(s_t, \mathbf{a}_{t,\tau}, \tau)$ from $\tau = 0$ to $\tau = 1$. The simplest integration method is Euler integration, which is given by the following update rule:

$$\mathbf{a}_{t,\tau+\frac{1}{n}} = \mathbf{a}_{t,\tau} + \frac{1}{n} \cdot v_\theta(s_t, \mathbf{a}_{t,\tau}, \tau), \quad (3)$$

which is repeated n times from $\tau = 0$ to $\tau = 1$ to obtain $\mathbf{a}_{t,1}$, where n is the number of integration steps (also called “denoising steps”). $\mathbf{a}_{t,1} = \mathbf{a}_t$ is the final action chunk that is executed as before.

1. **Implementation.** For this part, you will implement the `FlowMatchingSchedule` class in `networks.py` and `flow_matching_loss` in `losses.py`. Specifically, you should implement the following:

- `FlowMatchingSchedule.interpolate`: given a clean action chunk $\mathbf{a}_1$ and timestep $\tau \in [0, 1]$, sample noise and return the interpolated $\mathbf{a}_\tau$ and target velocity $v = \mathbf{a}_1 - \mathbf{a}_\tau$.
- `FlowMatchingSchedule.sample`: starting from $\mathbf{a}_0 \sim \mathcal{N}(0, I)$, run Euler integration for `num_steps` steps and return the result clamped to $[0, 1]$.
- `flow_matching_loss`: implement the flow matching loss presented above (tip: call `schedule.interpolate`)

2. **Run flow matching on hard mode** and report your results in a table (mean and standard deviation of episode length) [1 point]:

```
python main.py --method bc_flow --env hard
```

Solution: We expect a mean larger than 900

3. **Explain** in 2–3 sentences why flow matching performs the way it does on hard mode. [1 point]

Solution: The data on the hard mode is multimodal because of the double pipe opening. Flow matching is well suited to handle multimodality, and hence can learn to properly solve this environment.

Problem 3: DAgger [2 points]

In this section, you will implement the DAgger algorithm to improve the behavioral cloning (BC) policy trained with the regression objective in Problem 1.

As discussed in class, DAgger works by iteratively rolling out the current learned policy, π , to collect states encountered under its own behavior, which we store in $\mathcal{D}_\pi$. For each visited state s , the policy produces an action $\pi(s)$, but instead of keeping that action, we query the expert and relabel the state with the expert action. This gives the dataset

$$\mathcal{D}_{\text{DAgger}} = \{(s, \pi_{\text{expert}}(s)) \mid s \sim \mathcal{D}_\pi\}.$$

Once this relabeled dataset has been collected, we perform imitation learning on the aggregated dataset $\mathcal{D} \cup \mathcal{D}_{\text{DAgger}}$ using the same regression objective from Problem 1. This process is typically repeated for multiple DAgger iterations.

DAgger helps mitigate distribution shift between the expert and the learned policy by visiting states that might be present in the policy rollouts but not in the demos, and acquiring expert supervision in those states.

1. **Implementation.** Implement the following in `dagger.py`:

- `DeterministicExpert.act`: same logic as `Expert.act` in hard mode, but pick a strategy that will resolve the ambiguity found in the previous questions so that you make the MSE regression policy work.
- `rollout_episode`: roll out the current policy, don't forget to reset the environment and to use the action chunks from the policy and return the state action pairs.
- `rollout_and_relabel`: roll out the current policy using `rollout_episode`, use `DeterministicExpert` to relabel the episodes, and get the new action state training pairs.

2. **Run DAgger on hard mode** and report your results as a learning curve [1 point]:

```
python main.py --method dagger --env hard
```

Run dagger for 5 rounds (the default value). Plot the DAgger learning curve (round number on x-axis, mean episode length on y-axis, with error bars for standard deviation). Include the original regression performance as a horizontal line on the same plot.

Solution: We expect the mean of one of the iterations to be larger than 900

3. **Comparison.** Create a bar chart or table comparing the performance of all three methods on hard mode: Regression, Flow Matching, and DAgger (final round). [0.5 points]

```
python main.py --plot
```

Solution: +0.5 points for showing all of the methods on hard mode

4. **Explain** in 3–4 sentences: why does DAgger improve over rounds? What role does the deterministic expert play? How does this DAgger’s approach solve the challenges mentioned previously for the MSE regression policy? [0.5 points]

Solution: In our solution we choose to use DAgger to fix the MSE failing policy by removing the multimodality of the data and forcing the bird to always take the same hole of the pipe opening (either top or bottom). DAgger then iteratively improves the agent to learn to always take the same opening. Since this deterministic relabelling reduces the multimodality, now the MSE policy can successfully solve this task.

What do we consider a correct answer: Students should point out that they use DAgger to reduce the multimodality of the dataset by deterministically choosing one pipe opening [+0.25 points]. Students should say that through relabelling DAgger improves the policy and given the reduction in multimodality the MSE policy can now learn correctly [+0.25 points].