


RL for LLMs: Reasoning

Noam Brown
OpenAI

Scaling Laws

The incredible progress in AI from 2019 to today is primarily due to **scaling data and compute**



Is Scaling is All You Need?

You
I'll play in the bottom center.

ChatGPT
Nice move! Here's the updated board:

markdownCopy code

```
X | O |  
-----  
  | O |  
-----  
  | X | X
```

For my turn, I'll place my "O" in the top right:


markdownCopy code

```
X | O | O  
-----  
  | O |  
-----  
  | X | X
```

It's your turn again. Where would you like to place your next "X"?


Reasoning Models: A New Dimension for Scaling

- **Training** costs grew rapidly
- **Inference costs** were still low
 - Querying ChatGPT without reasoning cost <1 penny
- Reasoning models scale **inference** rather than just training




Annual Computer Poker Competition

- Each year, research labs would make poker bots and play them against each other.
- It turned into a competition of scaling models:




2015 Brains vs. AI Poker Competition

- In 2015 we (CMU) challenged 4 top poker pros to an 80,000-hand poker competition
- \$120,000 in prize money to incentivize them
- Our bot (Claudico) lost by 9.1 bb/100



The importance of planning in poker

Scaling in a medium-sized poker game [1]



[1] "Safe and Nested Subgame Solving in Imperfect-Information Games."
Brown & Sandholm. NeurIPS 2017 Best Paper.

2017 Brains vs AI Two-Player Poker AI

[Brown & Sandholm Science-17]

- Libratus against 4 top poker pros




- 120,000 hands of poker
- \$200,000 in prize money
- **Won by 15 bb/100** (Claudico had lost by 9 bb/100)
 - P-value ≈ 0.0002
- Each human lost individually to Libratus

Scaling Inference Compute in Backgammon

[Tesauro 1994]


- Achieved human master level in 1994
- First major success of neural networks in games
- **Used 2-3 ply search** in addition to its neural network



Scaling Inference Compute in Chess

[Campbell et al. 2002]

- Beat world champion Garry Kasparov in 1997
- Key to successes was large-scale alpha-beta pruning search
- Deep Blue spent minutes calculating each move




Scaling Inference Compute in Go

[Silver et al. Science-17]

Superhuman performance

No test-time search

- Increasing Elo by 120 points requires either:
 - ~2x model size and training
 - ~2x test-time search
- To get the raw policy net from 3000 Elo to 5200 Elo, you would need to scale by **~100,000x**



Is there a general way to scale
inference compute in LLMs?

Prompted Chain of Thought

Standard Prompting

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The answer is 27. 

Chain-of-Thought Prompting

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?


A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output


A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23 - 20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9. 

Prompted Chain of Thought



Verification in LLMs via Consensus / Majority Vote


- Generate a bunch of solutions and take the most common one
 - Minerva [Lewkowycz et al.] goes from 33.6% to 50.3% on MATH due to consensus
 - Consensus flatlines before 100 samples




Large Language Monkeys: Scaling Inference Compute
with Repeated Sampling

Bradley Brown^{*†‡}, Jordan Juravsky^{*†}, Ryan Ehrlich^{*†}, Ronald Clark[‡], Quoc V. Le[§],
Christopher Ré[†], and Azalia Mirhoseini^{†§}

OpenAI o1



OpenAI o1 → o3



Connections

March 5, 2025



Create four groups of four!

BITE	ONION	SOFT	KICK
PICKLE	SKIRT	ALLERGY	RELIEF
PAINT	DODGE	EDGE	CLIP
BORDER	FILE	MELODRAMA	FRINGE

Mistakes Remaining: ● ● ● ●

Solve this Connections puzzle:

BITE, ONION, SOFT, KICK, PICKLE, SKIRT, ALLERGY, RELIEF,
PAINT, DODGE, EDGE, CLIP, BORDER, FILE, MELODRAMA,
FRINGE

Reasoned about Connections puzzle for 1m 25s >

A concise way to see the four sets is:

1) BORDER, EDGE, FRINGE, SKIRT

All are synonyms for a boundary or perimeter.

2) BITE, CLIP, FILE, PAINT

All go with "nails": you can **bite** your nails, **clip** nails, **file** nails, and **paint** nails.


3) DODGE, KICK, PICKLE, SOFT

All can be followed by "ball": **dodgeball**, **kickball**, **pickleball**, and **softball**.

4) ONION, ALLERGY, RELIEF, MELODRAMA


All are associated with tears: onions make you cry, allergies cause watery eyes, "tears of relief," and a melodrama is a "tear-jerker."

GRPO algorithm




Group Relative Policy Optimization (GRPO) for LLM Reasoning


Optimize the policy using rewards relative to other samples from the same prompt.



R1-Zero: GRPO scaled up




DeepSeek-R1-Zero samples 16 outputs per question during RL



Accuracy and response length increase together during training

Chain of Thought vs Majority Voting




Claim: Need to rethink AI in an era of
large-scale test-time compute

GPT-5.5


	GPT-5.5	GPT-5.4	GPT-5.5 Pro	GPT-5.4 Pro	Claude Opus 4.7	Gemini 3.1 Pro
Terminal-Bench 2.0	82.7%	75.1%	-	-	69.4%	68.5%
Expert-SWE (Internal)	73.1%	68.5%	-	-	-	-
GDPval (wins or ties)	84.9%	83.0%	82.3%	82.0%	80.3%	67.3%
OSWorld-Verified	78.7%	75.0%	-	-	78.0%	-
Toolathlon	55.6%	54.6%	-	-	-	48.8%
BrowseComp	84.4%	82.7%	90.1%	89.3%	79.3%	85.9%
FrontierMath Tier 1–3	51.7%	47.6%	52.4%	50.0%	43.8%	36.9%
FrontierMath Tier 4	35.4%	27.1%	39.6%	38.0%	22.9%	16.7%
CyberGym	81.8%	79.0%	-	-	73.1%	-

GPT-5.5

CyberGym




Capture-the-Flags challenge tasks (Internal)




How far can we push test-time compute?

- o1 operated over seconds
- o3 operated over minutes
- Our IMO gold model operated over hours
- Today, scaffolds operate over days/weeks

How far can we push test-time compute?




Completed steps on "The Last Ones" per spent tokens




Multi-Agent AI

- CoT is inherently **serial**
 - Latency eventually becomes a bottleneck
- Some test-time scaling techniques are **parallel**
 - Best-of-N / consensus
 - Lower latency, but less compute-efficient




Multi-Agent AI


- CoT is inherently **serial**
 - Latency eventually becomes a bottleneck
- Some test-time scaling techniques are **parallel**
 - Best-of-N / consensus
 - Lower latency, but less compute-efficient



Multi-Agent AI




Score vs Compute/Time is the right way to eval



Score vs Compute/Time is the right way to eval

Intelligence vs. Output Tokens Used in Artificial Analysis Intelligence Index

Artificial Analysis Intelligence Index; Output Tokens Used in Artificial Analysis Intelligence Index





Safety/Preparedness evaluations are broken

- Safety evaluations measure whether a model might contribute to catastrophic harms (e.g., cyber, nuclear, biological weapons)
- These evaluations are typically done at relatively low budgets (<\$100)
- A dedicated state actor can easily spend \$10M on inference

Safety/Preparedness evaluations are broken

Gemini 3 Deep Think



**Zvi Mowshowitz** 
@TheZvi

[Subscribe](#)  

I confirmed with a Google representative that since this was a runtime improvement and they do not believe these performance gains constitute any additional risk, they believe that no safety explanation is required of them.

I found that to be a pretty terrible answer.

Methodology: deepmind.google/models/evals-methodology/gemini-3-deep-think

Safety evaluations should project test-time compute



Long-horizon safety evaluations are hard


- Suppose (for simplicity) we have 1T context and models operate over months
- How do we know what the behavior/capability will be after a month?
 - The only way to know for sure is to run it for a month

Inference capacity is strategically undervalued

- As inference becomes more important, weights become *relatively* less important
- Historically, big focus on securing model weights
 - Securing model weights still matters enormously!
- Inference compute is a strategic advantage

Test-time compute is a window into the future

- Capabilities that cost \$1M today might cost \$100 next year
- Large-scale test-time compute lets us anticipate future model capabilities
- We can use this time to prepare



Where does this go?

- There is still room to push inference compute **much further**
- Much **higher inference compute**, but much **more capable models**
 - What inference cost are you willing to pay for a proof of the Riemann Hypothesis?
 - What inference cost are you willing to pay for new life-saving drugs?
- Civilization was built by **billions** of humans over **millenia**
 - Similarly, there will likely be **billions** of persistent agents
 - In the same way that humans share knowledge, AI agents will share knowledge
 - In the same way humans specialize, AI agents will specialize

The Bitter Lesson by Richard Sutton

“The biggest lesson that can be read from 70 years of AI research is that general methods that leverage computation are ultimately the most effective... The two methods that seem to scale arbitrarily in this way are *search* and *learning*.”